

Scaling SSVD method and Command Line Interface notes

Dmitriy Lyubimov
dlieu.7@gmail.com

December 25, 2010

Contents

1	Base algoirthm	2
2	Matrix preprocessing	2
3	Enabling wider matrices.	4
4	Enabling taller matrices	4

1 Base algoirthm

Branch: *givens-ssvd*, options:

- k, --rank <int-value> the requested SVD rank (minimum number of singular values and dimensions in U, V matrices)
- p, --oversampling <int-value> stochastic SVD oversampling. (k+p=500 is probably more than reasonable).
- r, --blockHeight <int-value> the number of rows of source matrix for block computations. Taller blocking causes more memory use but produces less blocks and therefore somewhat better running times. However if blocks too tall, the algorithm may not be able to form them in the mapper in which case the split size may need to be increased with --minSplitSize (but overshoots over standard size are detrimental to network IO)
- s, --minSplitSize <int-value> the minimum split size to use in mappers. Since in this branch block formation happens in mappers, for significantly large -r and width of the input matrix the algorithm may not be able to read minimum k+p rows and form a block of minimum height, so the job would bail out at the very first mapping step. If that is the case, then one of the recourses available is to force increase in the MapReduce split size using SequenceFileInputFormat.setMinSplitSize() property. Increasing this significantly over HDFS size will result in network IO overhead as discussed in p.6.2 of the working notes).
- computeU <true|false> Request computation of the U matrix (default true) – computeV <true|false> Request computation of the V matrix (default true)
- reduceTasks <int-value> The number of reducers to use (where applicable): depends on size of the hadoop cluster

2 Matrix preprocessing

Branch: *ssvd-preprocessing*. *Ssvd-preprocessing* branch equips class *VectorWritable* with the following preprocessing interface.

The preprocessing interface *VectorPreprocessor* is introduced (alg. 1).

Algorithm 1 preprocessor interface

```
/**
 * Hooks into {@link VectorWritable} to provide matrix data preprocessing
 * capabilities to save on allocating long vectors .
 *
 * @author Dmitriy
 *
 */
public interface VectorPreprocessor extends Configurable {
/**
 * called before starting processing a new vector in the writable.
 *
 * @param sequential true if vector is either dense or sequential sparse.
 * false if vector formation is non-sequential.
 *
 * @return true if preprocessing of this vector type is supported.
 * if preprocessing is not supported, {@link VectorWritable} accumulates
 * the vector in the usual way instead of feeding to the preprocessor.
 */
boolean beginVector ( boolean sequential ) throws IOException;
/**
 * called when next vector element becomes available.
 * @param index
 * @param value
 */
void onElement ( int index, double value ) throws IOException;
/**
 * called for named vectors if vector name is available.
 *
 * @param name
 */
void onVectorName ( String name ) throws IOException ;
/**
 * called after last element is processed.
 */
void endVector () throws IOException;
}
```

The algorithmic change in Q-Job:

1. A row of \mathbf{A} is never formed. Instead, vector preprocessor is set up in the mapper and is accumulating $k + p$ dot products of incoming elements of \mathbf{A} . Thus, as soon as all elements of row \mathbf{A} are consumed, the preprocessor holds row of \mathbf{Y} . Preprocessor supports both sequential and non-sequential preprocessing.

The algorithmic change in \mathbf{B}^\top -Job:

1. A row of \mathbf{A} is never formed. Instead, vector preprocessor is set up in the mapper.
2. Partial \mathbf{B}^\top products $1 \times k + p$ are formed and output during preprocessing.
3. $\mathbf{Q}_{i,*}$ row is requested from preprocessor and $\hat{\mathbf{Q}}_j$ blocks are computed on the fly when current block is exhausted as a part of such request from preprocessor.

It resolves issue defined in § 6.2 of the “working notes” document.

3 Enabling wider matrices.

Branch: ssvd-wide. This branch enables processing wider matrices without having to increase `-minSplitSize` parameter. This is achieved by adding a reducer to Q-Job. The first Q pass is pushed down to reducers and mappers now only preprocess \mathbf{Y} rows. What it achieves is that mappers now will not fail if they can't acquire at least $k + p$ rows of \mathbf{A} . Most importantly, it defers “supersplits” problem until up to ~8 (16) mln dense items in incoming \mathbf{A} vectors assuming 64M(128M) HDFS blocks.

It alleviates problem defined in § 6.1 of the “working notes” document.

`-w, --wide <true|false>` enable first **QR** pass pushdown to reducers. **This will not be more efficient** unless matrix vectors significantly exceed ~10-50k non-zero elements.

4 Enabling taller matrices

It solves problem defined in § 6.3 of the “working notes” document.